

cmcR: Congruent Matching Cells method in R for cartridge case identification

by Joseph Zemmels, Heike Hofmann and Susan VanderPlas

Abstract Firearm evidence identification is the process of analyzing bullets or cartridge cases left at a crime scene to determine if they originated from a particular firearm. Statistical methods have long been developed and used to aid in such analyses. The Congruent Matching Cells (CMC) method is one such method developed at the National Institute of Standards and Technology (NIST) to quantify the similarity between two spent cartridge cases based on the markings left by the firearm barrel during the firing process. We introduce the first open-source implementation of the CMC method in the R package **cmcR**. The package will bolster forensic researchers' abilities to investigate, validate, and improve upon current statistical methodology in the field of forensic science.

Introduction

A *cartridge case* is a type of firearm ammunition that contains a projectile (e.g., bullet, shots, or slug). When a firearm is discharged, the projectile stored in the cartridge case is propelled down the barrel of the firearm. In response, the rest of the cartridge case that remains inside of the firearm is forced towards the back of the barrel. The force with which the cartridge case is propelled backwards causes it to strike the back wall, known as the *breech face*, of the barrel. Markings due to, e.g., manufacturing imperfections are ingrained on the breech face. When the cartridge case slams against the breech face, these markings can be "stamped" into either the primer of the cartridge case or the cartridge case itself. The markings left on a cartridge case from the firearm's breech face are called *breech face impressions*.

An example of the breech face from a 12 GAUGE, single-shot shotgun is shown in Figure 1a. The hole in the center of the breech face houses the firing pin that shoots out to strike a region on the base of the cartridge case known as the *primer*. This in turn ignites the propellant within the cartridge case causing a deflagration of gases that propels the bullet forward down the barrel. Figure 1b shows a cartridge case fired from the shotgun shown in Figure 1a. This cartridge case displays both a circular impression left by the firing pin in the middle of the primer as well as breech face impressions left on the outer region of the primer not impressed into by the firing pin.




Figure 1: Breech face of a barrel and associated breech face impressions on a cartridge case [Doyle (2019)].

These breech face impressions are considered to be analogous to a firearm's "fingerprint" left on a cartridge case. Matching an expended cartridge case of unknown source to one of known source based on breech face impressions has been performed for over 100 years by forensic practitioners [Thompson (2017)]. The development of computational and statistical methods to perform such identification has recently grown in interest [Council (2009)].

One such method is the Congruent Matching Cells (CMC) method developed at NIST that involves partitioning a cartridge case image or scan into a grid of "correlation cells" to isolate areas containing

identifying breech face impression markings [Song (2013)]. Since its invention in 2012, researchers at NIST have developed a number of extensions and improvements of the CMC method. However, to date there does not exist an openly available implementation of any of these techniques. Rather, many methods described in the CMC literature include a qualitative description of a proposed technique followed by results from the authors' implementation. The description of these methods do not delve into the intricacies of the implementation, which makes it especially difficult to validate or assess. Additionally, some procedures related to pre-processing the cartridge case data are seemingly done by-hand at NIST rather than with an automated method. This compounds the difficulty to accurately reproduce results. The **cmcR** package provides an open-source, fully-automatic implementation of the CMC method as originally described as well an extension proposed by Tong et al. (2015).

Cartridge case data

Cartridge case data commonly come in two forms: 2D grayscale images and 3D topographical scans. It is common in the CMC literature to use the 3D topographical scans to demonstrate the efficacy of a proposed method [Tong et al. (2015), Chen et al. (2017)]. A variety of scans are openly available for download through the NIST Ballistics Toolmark Research Database [Zheng et al. (2016)]. The **cmcR** package was designed specifically for use with the 3D topographies.

The 3D topographies are commonly stored in an .x3p (XML 3D Surface Profile) file format that includes metainformation such as who took the scan and the parameters under which the scan was taken (e.g., the lateral resolution in microns). The **x3ptools** package in R provides an interface to manipulate and visualize these .x3p files [Hofmann et al. (2019)]. The physical surface is represented using a *surface matrix*: a matrix of spatially-ordered elements or "pixels" whose values correspond to the height of the cartridge case surface at a particular location. Figure 2 shows the surface matrices of a known match (KM) pair of cartridge cases, meaning there were fired from the same firearm. Note that white regions in the images below represent unobserved or missing values. When read into R using the **x3ptools** package, these elements are encoded as NA. The size of a surface matrix depends on the lateral resolution with which the scans were taken. For example, a popular set of scans in the CMC literature were taken by Fadul et al. (2011); two of which are shown is the scan shown in Figure 2. These scans were taken with a lateral resolution of 6.25 microns per pixel. The actual surface matrices from this study vary around 1200×1200 pixels in size.




Figure 2: Two known match cartridge case scans from Fadul et al. (2011)

Only certain regions of a cartridge case contain identifying breech face impression markings. Song (2013) refers to these as "valid correlation regions" that are to be used to determine whether two cartridge cases match. The cell-based comparison procedure described in section 2.3.1 is designed to emphasize such regions. However, prior to applying this procedure cartridge scans must undergo some pre-processing to remove sections of the cartridge case surface that do not come into contact with the breech face of the barrel. These include a circular plateaued region in the center of the scan that is pushed aside by the firing pin during the firing process and clusters of observed values in the corners of the scan that are artifacts of the staging area in which the scan was captured. The task in pre-processing is to automatically remove these unwanted regions from the scan to accentuate unique markings left by the breech face. This is discussed in greater detail in section (LINK).

Cell-based surface matrix comparisons

This section will detail how two cartridge case scans can be compared and classified as a match using a cell-based procedure.

Cell-based comparison procedure

The Congruent Matching Cells method was developed at the National Institute of Standards and Technology to quantify the similarity between two spent cartridge cases based on their breech face impressions. The CMC method involves dividing a breech face impression scan into a grid of cells and comparing each cell in one scan to a corresponding region in the other scan. This method is motivated by the fact that breech face markings are not uniformly impressed upon the cartridge case during the firing process. As such, only certain sections of the cartridge case have identifiable markings that make it possible to match to a firearm. Calculating a similarity score between the entirety of two cartridge case surfaces might not highlight these identifying regions. Instead, the number of highly similar cell pairs between the two scans can be used as a more granular similarity metric.

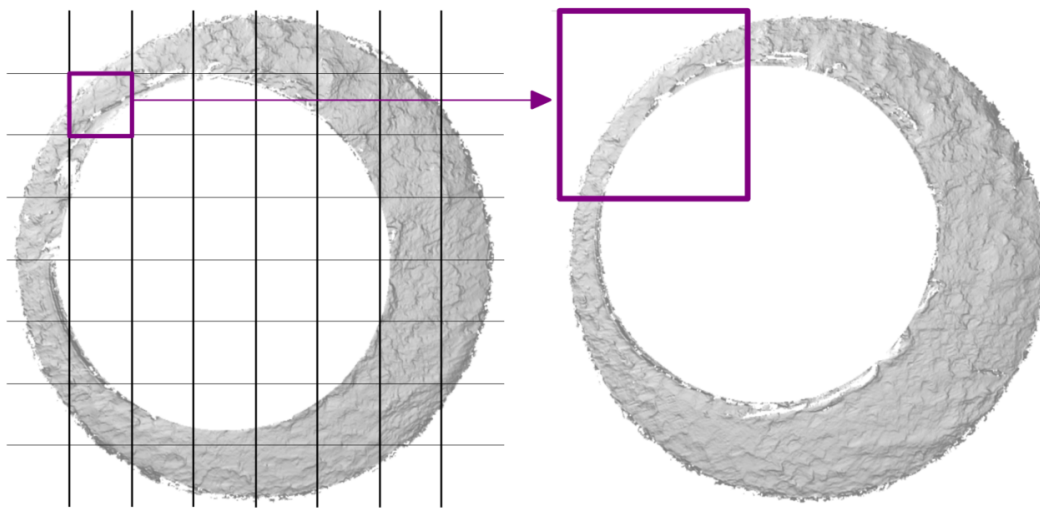


Figure 3: Illustration of comparing a “cell” in one cartridge case scan to a region in another.

Figure 3 illustrates the cell-based comparison procedure between two cartridge case scans. The scan on the left is divided into a grid of 8×8 cells. Each cell is paired with an associated larger region in the other scan. The absolute location of each cell and region in their respective surface matrices remain constant. However, the scan on the right is rotated to determine the rotation at which the two scans are the most “similar,” which is quantified using the *cross-correlation function* (CCF). For two real-valued, $M \times N$ matrices A and B , the cross-correlation function, denoted $(A \star B)$ can be defined as

$$(A \star B)[m, n] = \sum_{i=0}^M \sum_{j=0}^N A[i, j] B[(i + m)_{\text{mod } M}, (j + n)_{\text{mod } N}].$$

Note that this finite, discretized CCF is a matrix of elements representing the similarity between matrices A and B for various translations of matrix B . The index at which the CCF attains a maximum represents the translations needed to align B with A . In practice, calculating the CCF from the definition is often computationally intractable. The *Cross-Correlation Theorem* provides a feasible alternative to calculating the CCF. For two matrices A and B , the Cross-Correlation Theorem says that

$$(A \star B)[m, n] = \mathcal{F}^{-1} \left(\overline{\mathcal{F}(A)} \cdot \mathcal{F}(B) \right) [m, n]$$

where \mathcal{F} and \mathcal{F}^{-1} denote the discrete Fourier and inverse discrete Fourier transforms, respectively [Brigham (1988)]. Note that the multiplication on the right-hand side is pointwise (Hadamard) multiplication. This result allows us to trade the moving sum computations from the definition of the CCF for two forward Fourier transformation, a pointwise product, and an inverse Fourier transformation. The Fast Fourier Transform (FFT) algorithm is used to reduce the computational load considerably. However, a practical consideration for applying this method with cartridge case data is the large number of non-random missing values in a surface matrix. Recall that missing values are represented in Figures 2 and 3 as white pixels. The discrete Fourier transform is not defined for

matrices containing missing values, so these need to be replaced. The convention adopted in the **cmcR** package is to replace missing values with 0 after standardizing a matrix by subtracting away its average height value and dividing by its standard deviation. Such standardization is commonly performed by authors at NIST [for example, Ott et al. (2017)]. While replacing missing values is essential for using the FFT-based method of calculating the CCF, doing so causes the CCF values to be “deflated” relative to the pairwise-complete cross-correlation in which only pairs of pixels in which neither element is missing are considered. However, the translation estimates obtained from this method are often good estimates for true translation values by which the two matrices align.

Figure 4 provides an example of the output from the FFT-based CCF calculation method. In the top-left we see a 72×72 pixel cell from one surface matrix. In the top-right we see this cell’s associated region in the other surface matrix of dimension 216×216 (triple the cell’s side lengths). The bottom-left shows the CCF “map” calculated using this FFT-based method. Although the CCF need not be bounded between -1 and 1 based on the definition, it is common to normalize the CCF for interpretability purposes and is done so in the **cmcR** package. A summary of the alignment parameters at which the CCF_{\max} occurs is shown in the bottom-right. We can see that the two matrices are best-aligned when the cell is shifted “east” 20 pixels and “north” 9 pixels starting from the center of the region. The $\theta = -18$ indicates that the overall cartridge case scan from which the region (shown in the top-right) was extracted was rotated by -18 degrees for this comparison. The orange square in the top-right plot shows where the cell would be located if the translation were performed.

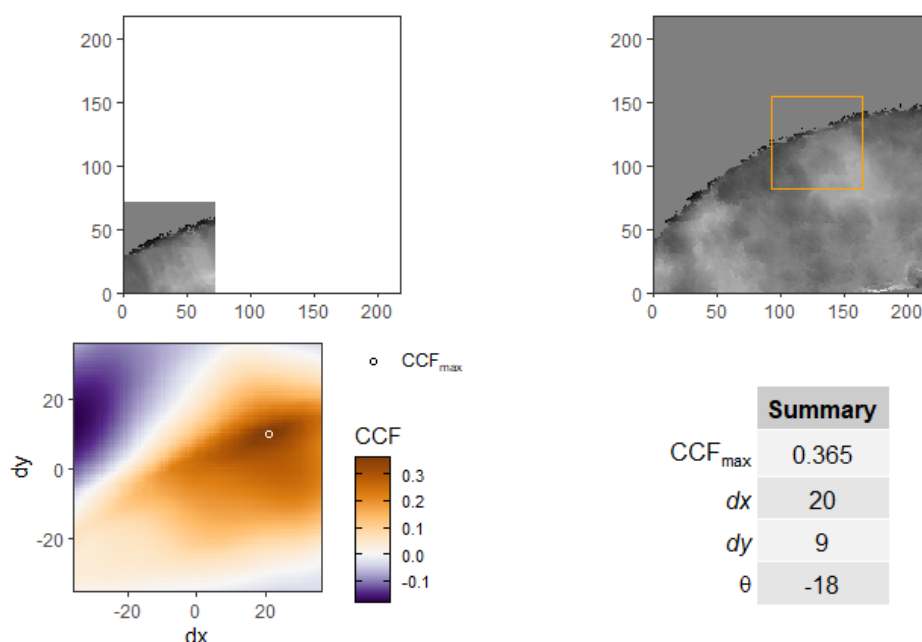


Figure 4: Example of a cross-correlation function “map” for a particular cell/region comparison.

Using the estimated translation values at which the CCF_{\max} occurs, we can calculate the pairwise-complete cross-correlation between the cell and a cell-sized matrix extracted from the larger region where missing values are not replaced. Think of this as punching-out the matrix enclosed in the orange square shown in the top-right plot of Figure 4. This will be used as the final CCF_{\max} estimate. This cell-based comparison procedure is performed for each cell/region pair for various rotation values.

The Congruent Matching Cells method

A particular cell/region pair is deemed “highly similar” if it passes a collection of user-defined similarity criteria. The criteria are based on the fact that a pair of matching cartridge case scans are not necessarily aligned properly in their raw format. In particular, one cartridge case scan needs to be rotated and translated to align correctly the other. These unknown alignment parameters can be estimated for each cell/region pair using the cell-based comparison procedure discussed in section 2.3.1. For a truly matching pair of cartridge cases, we would expect these alignment parameter estimates to agree with each other across cell/region pairs; at least up to some threshold. Conversely, we would expect the estimates to vary randomly for a truly *non*-matching pair. As such, the CMC method details how to determine whether a consensus exists among the estimated alignment parameter values across the cell/region pairs.

Since the initial proposal of the CMC method, a number of improvements and extensions have been discussed by various authors. The **cmcR** package provides an implementation for the initially proposed method as well one extension proposed by [Tong et al. \(2015\)](#).

The initially proposed method

In their raw format, a pair of matching cartridge case scans are not necessarily properly aligned. If the surface matrices of two scans were stacked on top of one another such that their centers coincided, one would likely need to be rotated and translated to achieve maximum similarity with the other. Once this is accomplished, an accurate similarity score can be calculated between the two scans. For two scans A and B , let $\beta = (dx, dy, \theta)' \in \mathbb{R} \times \mathbb{R} \times [0, 2\pi)$ denote the translations and rotation needed to properly align A to B (so $-\beta$ would align B to A). For the CMC method, we divide A into n cells and pair each cell with a region in B . As illustrated in Figure ??, the regions in B are defined based on the location of their associated cell in A , yet with larger side lengths.

Each cell will have an associated $\beta_i = (dx_i, dy_i, \theta_i)', i = 1, \dots, n$, by which it will achieve maximum similarity to the corresponding region in B . If A and B are truly matching (were fired from the same firearm), then we would expect $\beta_i = \beta$ for each i . Due to the unknowable nature of β , the best we can hope to determine is if $\beta_i = \beta_j, i \neq j$. That is, one criterion specified to call a pair of cartridge case scans a match is that their cells' alignment parameters must agree. For non-matches, we would not expect the $\{\beta_i\}_{i=1, \dots, n}$ to agree in this way but rather to differ randomly. Note that, as shown in Figure ??, some cells contain little to no observed values; for example, the corner cells of the Image 1 scan. Such cells are not considered when determining similarity, although some vaguity exists in how one might define a threshold to classify a particular cell as containing too few observed values. For example, Chen et al. (2017) (SOURCE) only consider cells containing 15% or more observed values.

Because the cartridge cases are represented digitally, it is impossible to search over every possible value of β_i . Instead, the alignment criterion, $\beta_i = \beta_j, i \neq j$, is applied to estimates for the alignment parameters, $\{\hat{\beta}_i\}_{i=1, \dots, n}$. The translation parameter estimates are determined using the maximum of the cross-correlation function (CCF). This is effectively a template-matching procedure, borrowing terminology from image processing, using the CCF as a similarity metric. The rotation parameter estimate is determined through a grid search. The method provides some estimation leeway by specifying a threshold, $\mathbf{T} = (T_{dx}, T_{dy}, T_{\theta})'$, within which we say that two $\hat{\beta}_i$ s agree. The CCF is also used to quantify similarity between two cells analogous to how correlation can be used to measure the linear relationship between two random variables. In fact, a normalized CCF bounded between -1 and 1 is commonly considered for interpretability. An additional criterion to classify a cell pair as matching is a minimum CCF threshold, T_{CCF} . Note that these thresholds differ between authors (SOURCES).

In summary, the rotation parameter θ is estimated by comparing scan A to scan B over a grid of rotations of scan B . For each rotation of B , each cell in A is compared to its associated region in B using the CCF. This comparison yields both a maximum cross-correlation function value and an estimate of $(dx, dy)'$ for that particular rotation value. The method as initially proposed keeps track of the estimated alignment parameters by which each cell in scan A achieves its maximum CCF with the associated region in scan B . That is, each cell gets a single "vote" for the $\hat{\beta}_i$ that provides the best alignment between it and its associated region in scan B . Cells that contain too few observed values aren't given a vote. A "consensus" is determined by aggregating in some way the $(dx, dy, \theta)'$ votes. For example, the median of the $(dx, dy, \theta)'$ estimates may be taken as the consensus. The votes of each cell are then compared to this consensus to determine if they are within the $(T_{dx}, T_{dy}, T_{\theta})'$ thresholds. Any cell whose $\hat{\beta}_i$ vote fall within these thresholds and whose maximum CCF is above the specified T_{CCF} threshold is classified as a "congruent matching cell."

The High CMC method

[Tong et al. \(2015\)](#) discuss an improved version of the CMC method that targets specific deficiencies in the initially proposed method. Namely, the fact that each image A cell gets only one vote for the $(dx, dy, \theta)'$ alignment parameters by which it achieves maximum similarity with its associated region in image B . A cell pair may be wrongfully excluded from the "congruent matching cell" classification because, for example, the rotation for which it votes falls too far from the consensual rotation value. In reality, the cell pair may be highly similar at the consensual rotation value, yet is never given the chance to demonstrate this. The improved version of the method attempts to effectively give such cell pairs a second chance by considering their behavior at the consensual rotation value; although it introduces an additional criterion to guard against the possibility of a non-matching cell pair being highly similar by chance at the consensual rotation value.

This additional criterion is based on an observation that appears to be true based on empirical

evidence: if a pair of cartridge case scans A and B are truly matching after rotating B by θ , then they should still be highly similar after rotating B by $\theta + \epsilon$ for some small ϵ . In terms of the CMC method, this means that the number of congruent matching cell pairs should reach a mode at the true θ value. In practice, because only a grid of θ values are considered, a mode is detected based on how close the θ values yielding a high number of CMCs are to each other. Tong et al. classifies a “high” number of CMCs as $\text{CMC}_{\max} - \tau$ where CMC_{\max} is the maximum number of CMCs detected across all θ values and τ is an empirical constant ($\tau = 1$ in Tong et al. (2015)).

Another improvement to the original method proposed by Tong et al. is to consider both a comparison of scan A to rotated versions of scan B as well as B to A . The motivation for this being that true matches should be highly similar in either direction while non-matches should have conflicting or random results. Figure 5 shows the “forward” and “backward” CMC count distribution for a known match pair across a grid of rotation values, $\theta \in \{-30, -27, \dots, 27, 30\}$. We can see that both distributions achieve a mode at opposite θ values, which is expected and desired for a known match pair. Additionally, a horizontal line has been drawn at $\text{CMC}_{\text{high}} = \text{CMC}_{\max} - 1$ for the respective distributions. We can see that in either case, any θ value that achieves this CMC_{high} value is close to the mode. In contrast, Figure 6 shows the CMC count distribution for a known non-match pair that is relatively flat. In particular, we see that there are θ values achieving CMC_{high} that are far away from the mode.




Figure 5: CMC count per rotation (θ) value for Forward (A vs. B) and Backward (B vs. A) comparison of a known match cartridge case scan pair




Figure 6: CMC count per rotation (θ) value for Forward (A vs. B) and Backward (B vs. A) comparison of a known non-match cartridge case scan pair

These additional “high CMC count” criteria provide more enhanced discriminatory ability between known matching and known non-matching scans. If this criterion is satisfied by a particular pair of scans and a θ mode is identified, then the method considers CCF values for each cell pair at and around this mode. Again, this is in contrast to the initially proposed method that only considers the CCF_{\max} value and resulting (dx, dy, θ) votes for each cell pair.

The cmcR package

This section will highlight the **cmcR** package's functionality by walking through a possible use case. Many of the functions in this package provide the user with a variety of processing options with which they can experiment. This is due to the fact that processing techniques differ considerably among authors or are not discussed in great detail.

Pre-processing procedures

Studies in which cartridge cases are matched by forensic examiners often involve giving examiners a set of known matches and asking them to classify additional matches from a collection of unknown source scans. For brevity, we will consider a comparison between two cartridge cases. This particular pair of scans, as well as many other cartridge case scans, are openly available from the NIST Ballistics and Toolmarks Research Database. Figure 7 shows the pair of scans after performing the necessary pre-processing procedures. Note that the color scheme has been scaled by quantiles to visually highlight regions of the cartridge case scan containing strong breech face impression "signal."

The family of functions in the **cmcR** package beginning with `preProcess_` can be used to perform the necessary pre-processing steps for a pair of cartridge case scans to be comparable using the cell-based comparison procedure outlined in section 2.3.1. The implementation of many of these pre-processing procedures is inspired largely by [Tai and Eddy \(2018\)](#) who detail a fully-automatic procedure for processing cartridge case 2D images as opposed to 3D scans. The functions available include:

1. `preProcess_ransac`: estimates the height value of the breech face impressions in a scan using the Random Sample Consensus (RANSAC) robust, iterative plane-fitting algorithm [[Fischler and Bolles \(1981\)](#)].
2. `preProcess_levelBF`: extracts the observations containing breech face impressions from the scan using the estimated height value obtained from `preProcess_levelBF`.
3. `preProcess_cropWS`: removes rows/columns containing mostly if not all NA values from the surface matrix on the exterior of breech face impressions.
4. `preProcess_removeFPCircle`: detects and removes observations within the firing pin impression circle using the Hough Transform circle detection algorithm [[Hough \(1962\)](#)].
5. `preProcess_gaussFilter`: applies a low-pass, high-pass, or band-pass Gaussian filter to the breech face impressions to reduce the effects of high frequency noise, low frequency global structure, or both, respectively.

See the **cmcR** package documentation for more information about these functions.

For computational purposes it is common the CMC literature to down-sample a surface matrix prior to performing the cell-based comparison procedure. The `sample_x3p` function from the **x3ptools** package can be used to sample every m th row/column of a surface matrix. Together with the `preProcess_` family of functions, the code to produce the first surface matrix shown in Figure 7 is given by the following example.

```
library(cmcR)
library(x3ptools)
library(magrittr)

nrbsd_url <- "https://tsapps.nist.gov/NRBD/Studies/CartridgeMeasurement/"

fadul1.1_id <- "DownloadMeasurement/2d9cc51f-6f66-40a0-973a-a9292dbec36d"
fadul1.2_id <- "DownloadMeasurement/cb296c98-39f5-46eb-abff-320a2f5568e8"

fadul1.1 <- read_x3p(file = paste0(nrbsd_url,fadul1.1_id)) %>%
  sample_x3p(m = 2)

fadul1.1$surface.matrix <- fadul1.1$surface.matrix %>%
  #first RANSAC estimation yields rough height estimate
  preProcess_ransac(inlierThreshold = 1e-6, #1 micrometer
                    iters = 150) %>%
  preProcess_levelBF() %>%
  #second, more precise RANSAC estimation
  preProcess_ransac(inlierThreshold = 1e-7,
                    iters = 300) %>%
```

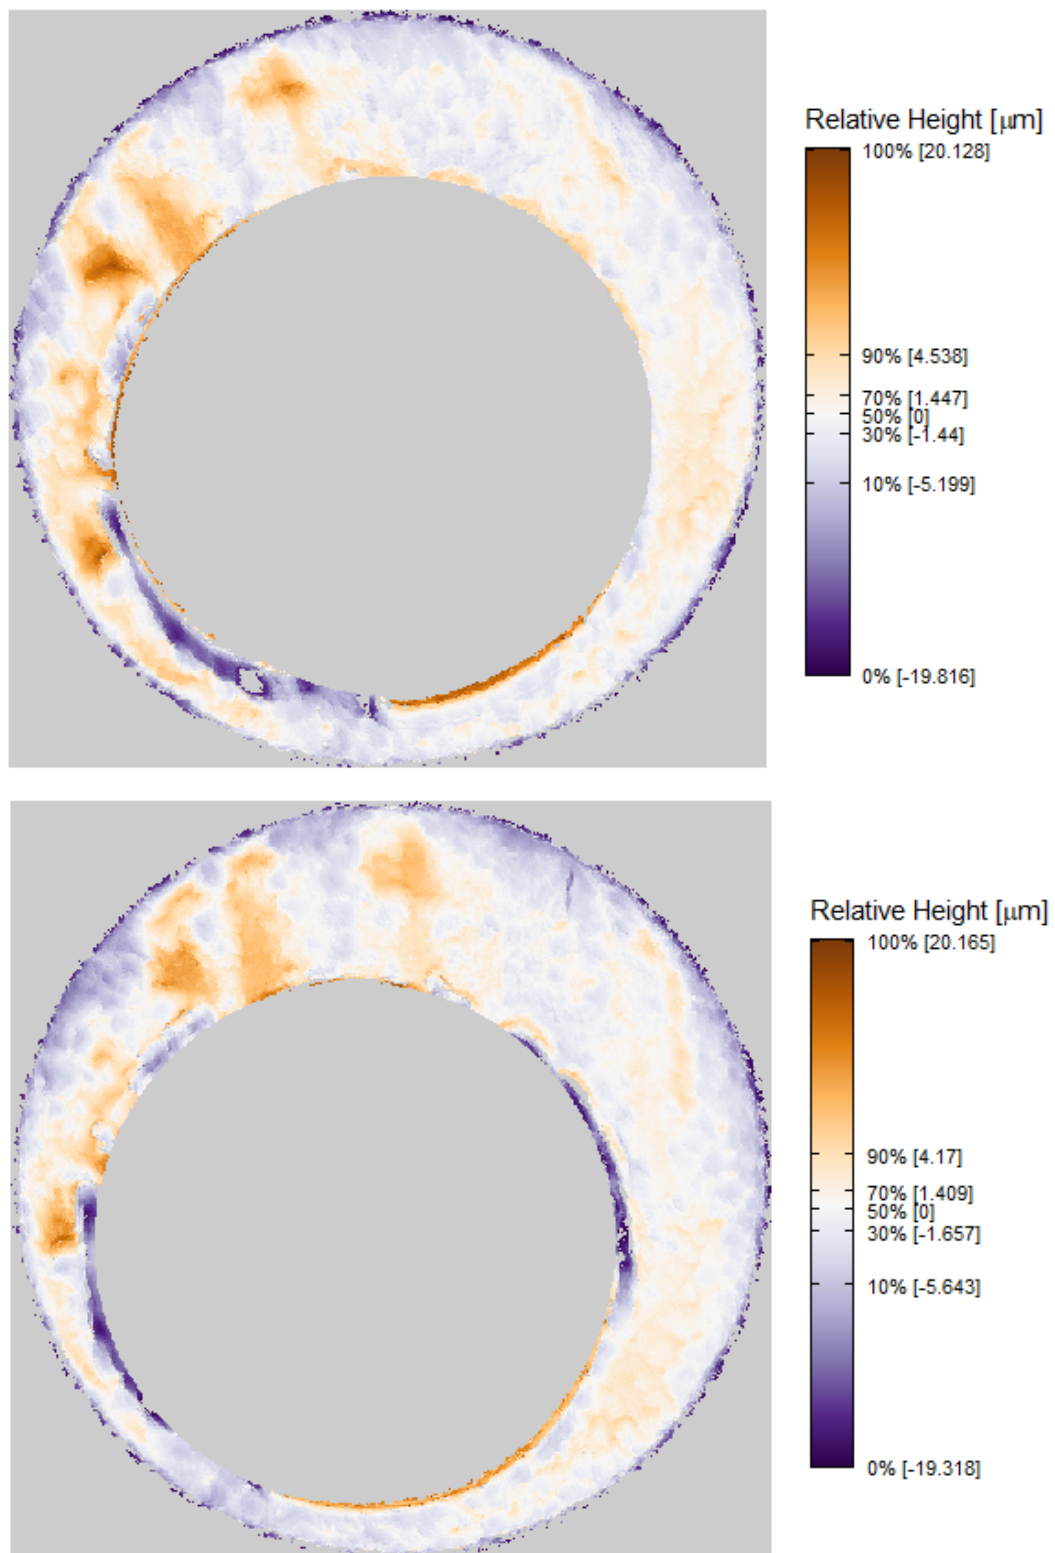


Figure 7: A known match pair of processed cartridge case scans.


```
preProcess_levelBF() %>%
preProcess_cropWS(croppingThresh = 1) %>%
preProcess_removeFPCircle() %>%
#Gaussian band-pass filter
preProcess_gaussFilter(res = fadul1.1$header$info$incrementY,
                        wavelength = c(16,250),
                        filtertype = "bp")
```

The same procedure can be performed to produce the second surface matrix shown in Figure 7, fadul1.2.

Implementation of cell-based comparison procedure

The cell-based comparison procedure outlined in section 2.3.1 is implemented in the cellCCF_bothDirections function. In particular, the procedure is performed twice so that both cartridge case scans take-on the role of the scan that is partitioned into a grid of cells. This is necessary to apply the High CMC logic discussed in section 2.3.4. Continuing with the current use case example, the code to perform this procedure on fadul1.1 and fadul1.2 is given by the following example.

```
kmComparison <- cellCCF(x3p1 = fadul1.1,
                        x3p2 = fadul1.2,
                        thetas = seq(-30,30,by = 3),
                        cellNumHoriz = 8,
                        cellNumVert = cellNumHoriz)
```

[Put cellCCF table here]

Congruent Matching Cells logic

With the CCF results calculated across a range of rotation values, we can extract the features used in the Congruent Matching Cells method.

The `cmcR::topResultsPerCell` function extracts from the list returned by `cellCCF` information about the $(\theta, dx, dy)'$ values at which each cell pair achieves its maximum CCF. The resulting data frame is required for determining the number of CMCs under the initially proposed method. Recall that the initially proposed method filters this data frame down by determining which cell pairs have associated $(\theta, dx, dy)'$ values close to the consensus $(\theta, dx, dy)'$. Passing this data frame to the `cmcR::cmcFilter` function implements this filtering process. The `cmcFilter` function provides options to specify the function used to form a consensus among $(\theta, dx, dy)'$ values and define thresholds for calling a particular cell pair's $(\theta, dx, dy)'$ values "close" to the consensus. Optionally, separate consensus functions can be defined for the θ and $(dx, dy)'$ values. This is motivated by the fact that, especially for a coarse grid, the θ values can be considered categorical while the $(dx, dy)'$, although only ever observed as integers, should be treated as continuous. The code below illustrates how CMCs under the initially proposed method can be calculated. The `cmcR::getMode` function calculates the mode θ value among those in the `topResultsPerCell` data frame.

```
comparison1$ccfResults %>%
topResultsPerCell() %>%
cmcFilter(consensus_function = median,
          ccf_thresh = .55,
          dx_thresh = 35,
          dy_thresh = 35,
          theta_thresh = 3,
          consensus_function_theta = cmcR::getMode)
```

Table (REF) shows the CMCs determined by the `cmcFilter` function for the two comparisons made in the current use case. We can see that the comparison between Unknown Source and Known Source 1 yields more CMCs than the comparison with Known Source 2. In fact, Song et al. (2013) (CHECK THIS SOURCE) recommend a minimum of 6 CMCs before calling a cartridge case pair a match. For this fabricated use case, Unknown Source indeed matches to Known Source 1 while Known Source 2 is a known non-match.

While useful for an initial estimate of the CMC count between two cartridge case scans, the improved method as described in section (LINK) has been shown to be more effective at classifying

cellNum	cellID	ccf	dx	dy	θ
5	y = 1 - 73, x = 289 - 360	0.56	-1	3	18
15	y = 74 - 145, x = 433 - 504	0.46	11	13	21
39	y = 291 - 362, x = 433 - 504	0.43	-1	-1	18
46	y = 363 - 435, x = 361 - 432	0.45	1	0	18
50	y = 436 - 507, x = 73 - 144	0.58	-6	-7	21
59	y = 508 - 579, x = 145 - 216	0.41	1	0	18

Table 1: 6 initial CMCs for the comparison under consideration.

matches from non-matches (SOURCES AGAIN?). The `cmcR::cellCCF_bothDirections` is just a wrapper for calling `cellCCF` twice where the two `x3p` objects change roles for the second call. The value returned by `cellCCF_bothDirections` is a list of two elements, the first being the comparisons between `x3p1` to rotations of `x3p2` and the second being the comparisons between `x3p2` to rotations of `x3p1`. The output of `cellCCF_bothDirections` can be handed off to the `cmcR::cmcFilter_improved` function to implement the improved CMC method. The value returned by `cmcR::cmcFilter_improved` is a list of 3 elements. The first element, `params`, contains the argument specifications under which the function call was made. The second element, `initialCMCs`, contains two dataframes of the CMCs calculated under the initially proposed method (equivalent to calling `cmcFilter` twice). The third element, `finalCMCs`, contains a data frame with the CMCs calculated under the improved method. If no additional CMCs are determined using the improved method, then a cartridge case pair is assigned the number of CMCs calculated under the initially proposed method. Such is the case for the comparison between Unknown Source and Known Source 2. Table (REF) shows this table for the current use case between Unknown Source and Known Source 1.

```
comparison2 <- cmcR::cellCCF_bothDirections(x3p1 = processedBF1$x3p,
                                           x3p2 = processedBF2$x3p,
                                           thetas = seq(-30,30,by = 3),
                                           cellNumHoriz = 8,
                                           cellNumVert = 8,
                                           centerCell = "wholeMatrix")

cmcR::cmcFilter_improved(cellCCF_bothDirections_output = comparison2,
                         consensus_function = median,
                         ccf_thresh = .5,
                         dx_thresh = 35,
                         dy_thresh = 35,
                         theta_thresh = 3,
                         consensus_function_theta = cmcR::getMode)
```

The final CMCs are shown in Figure (REF).

[Put plot of selected CMCs here]

Discuss here the CMC distribution for all known match and known non-match pairs?

Results

Implementation accuracy

Implementation sensitivity

Conclusion

Bibliography

- E. O. Brigham. *The Fast Fourier Transform and Its Applications*. Prentice-Hall, Inc., USA, 1988. ISBN 0133075052. [p3]
- Z. Chen, J. Song, W. Chu, J. A. Soons, and X. Zhao. A convergence algorithm for correlation of breech face images based on the congruent matching cells (CMC) method. *Forensic Science International*, 280:

cellNum	cellID	ccl	dx	dy	θ	comparison
3	y = 1 - 73, x = 146 - 217	0.52	13	-10	-15	comparison_1to2
4	y = 1 - 73, x = 217 - 288	0.50	-2	2	21	comparison_2to1
5	y = 1 - 73, x = 289 - 360	0.56	-1	3	18	comparison_2to1
10	y = 74 - 145, x = 73 - 144	0.58	7	-3	21	comparison_2to1
14	y = 74 - 145, x = 362 - 433	0.43	7	6	-21	comparison_1to2
15	y = 74 - 145, x = 433 - 504	0.46	0	3	18	comparison_2to1
17	y = 146 - 218, x = 1 - 72	0.44	-1	6	21	comparison_2to1
18	y = 146 - 218, x = 74 - 145	0.46	-6	6	-21	comparison_1to2
23	y = 146 - 218, x = 434 - 505	0.42	-1	-3	-18	comparison_1to2
24	y = 146 - 218, x = 506 - 577	0.44	1	1	-15	comparison_1to2
25	y = 219 - 290, x = 1 - 72	0.64	-1	4	18	comparison_2to1
31	y = 219 - 290, x = 433 - 504	0.44	-5	8	15	comparison_2to1
33	y = 291 - 362, x = 1 - 72	0.45	0	5	18	comparison_2to1
39	y = 291 - 362, x = 434 - 505	0.53	-1	-1	-21	comparison_1to2
41	y = 363 - 435, x = 1 - 73	0.56	-10	-12	-15	comparison_1to2
42	y = 363 - 435, x = 74 - 145	0.50	7	6	-21	comparison_1to2
46	y = 363 - 435, x = 361 - 432	0.45	1	0	18	comparison_2to1
48	y = 363 - 435, x = 505 - 576	0.50	4	-10	21	comparison_2to1
50	y = 436 - 507, x = 73 - 144	0.58	-6	-7	21	comparison_2to1
53	y = 436 - 507, x = 290 - 361	0.51	8	-5	-21	comparison_1to2
55	y = 436 - 507, x = 434 - 505	0.46	0	-1	-18	comparison_1to2
59	y = 508 - 579, x = 146 - 217	0.55	-6	0	-18	comparison_1to2
60	y = 508 - 579, x = 217 - 288	0.61	10	1	15	comparison_2to1
61	y = 508 - 579, x = 289 - 360	0.55	7	-2	18	comparison_2to1
62	y = 508 - 579, x = 362 - 433	0.42	-9	4	-15	comparison_1to2

Table 2: 25 final CMCs for the comparison under consideration.




Figure 8: Separated KM and KNM CMC count distributions under the initially proposed method.




Figure 9: Separated KM and KNM CMC count distributions under the improved method.

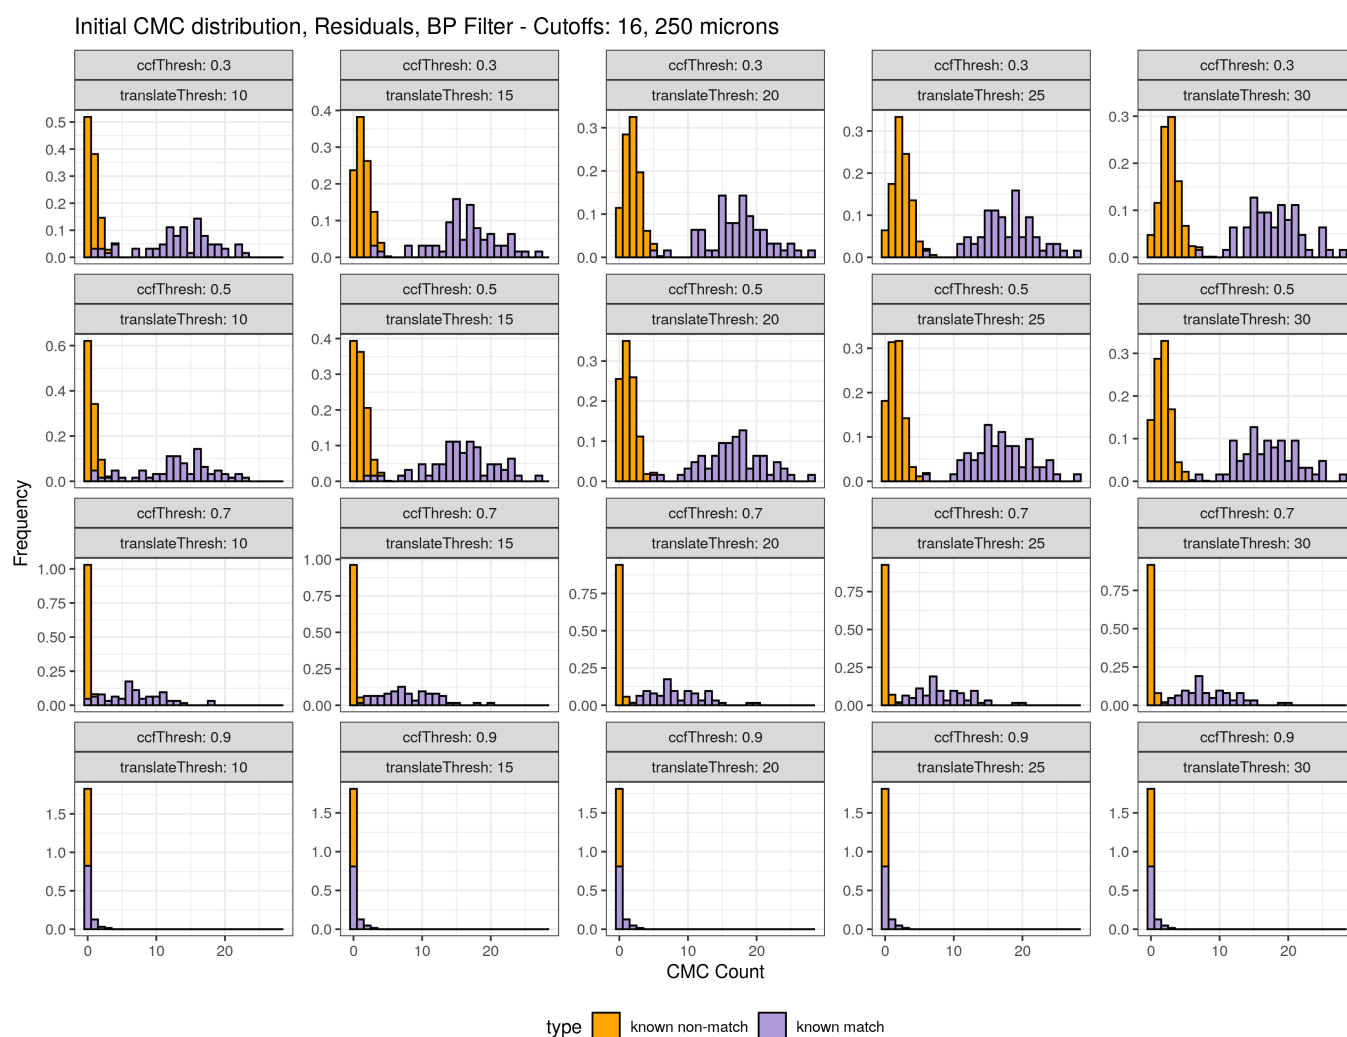


Figure 10: KM and KNM CMC count distributions under the initial method for various combinations of CCF_{\max} and translation thresholds.

CMC distribution, Residuals, BP Filter - Cutoffs: 16, 250 microns

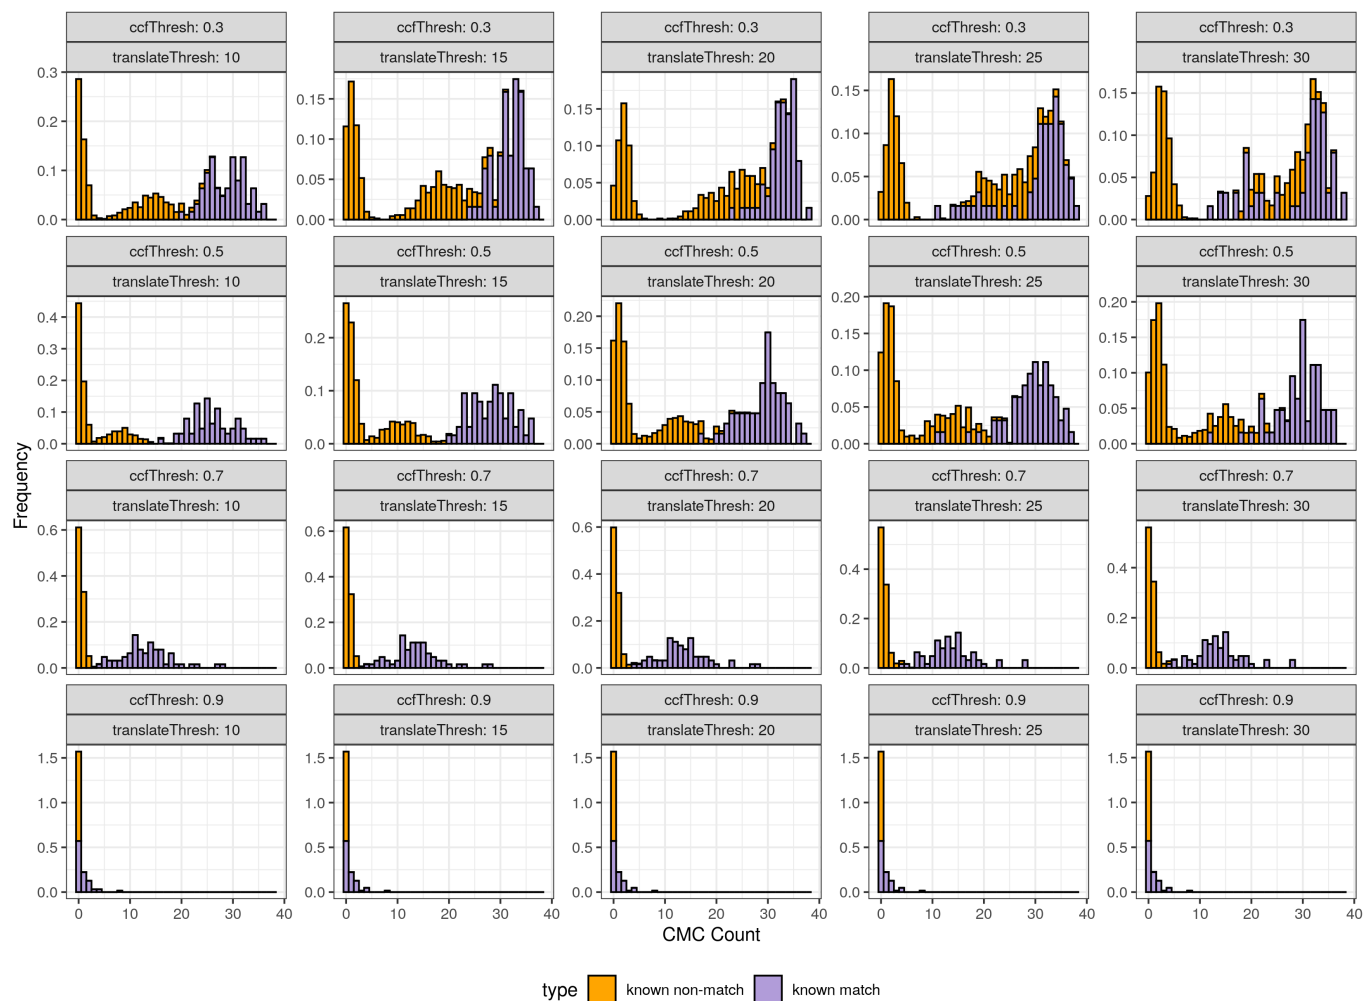


Figure 11: KM and KNM CMC count distributions under the improved method for various combinations of CCF_{max} and translation thresholds.

- 213–223, Nov. 2017. ISSN 03790738. doi: 10.1016/j.forsciint.2017.08.033. URL <https://linkinghub.elsevier.com/retrieve/pii/S0379073817303420>. [p2]
- N. R. Council. *Strengthening Forensic Science in the United States: A Path Forward*. The National Academies Press, Washington, DC, 2009. ISBN 978-0-309-13130-8. doi: 10.17226/12589. URL <https://www.nap.edu/catalog/12589/strengthening-forensic-science-in-the-united-states-a-path-forward>. [p1]
- J. S. Doyle. Cartridge case identification, 2019. URL http://www.firearmsid.com/A_CCIDImpres.htm. [p1]
- T. Fadul, G. Hernandez, S. Stoiloff, and G. Sneh. An Empirical Study to Improve the Scientific Foundation of Forensic Firearm and Tool Mark Identification Utilizing 10 Consecutively Manufactured Slides, 2011. [p2]
- M. A. Fischler and R. C. Bolles. Random sample consensus: A paradigm for model fitting with applications to image analysis and automated cartography. *Commun. ACM*, 24(6):381–395, June 1981. ISSN 0001-0782. doi: 10.1145/358669.358692. URL <https://doi.org/10.1145/358669.358692>. [p7]
- H. Hofmann, S. Vanderplas, G. Krishnan, and E. Hare. *x3ptools: Tools for Working with 3D Surface Measurements*, 2019. URL <https://github.com/heike/x3ptools>. R package version 0.0.2.9000. [p2]
- P. Hough. Method and means for recognizing complex patterns, 1962. [p7]
- D. Ott, R. Thompson, and J. Song. Applying 3D measurements and computer matching algorithms to two firearm examination proficiency tests. *Forensic Science International*, 271:98–106, Feb. 2017. ISSN 03790738. doi: 10.1016/j.forsciint.2016.12.014. URL <https://linkinghub.elsevier.com/retrieve/pii/S0379073816305461>. [p4]
- J. Song. Proposed “NIST Ballistics Identification System (NBIS)” Based on 3D Topography Measurements on Correlation Cells. *American Firearm and Tool Mark Examiners Journal*, 45(2):11, 2013. [p2]
- X. H. Tai and W. F. Eddy. A Fully Automatic Method for Comparing Cartridge Case Images,. *Journal of Forensic Sciences*, 63(2):440–448, Mar. 2018. ISSN 00221198. doi: 10.1111/1556-4029.13577. URL <http://doi.wiley.com/10.1111/1556-4029.13577>. [p7]
- R. Thompson. Firearm identification in the forensic science laboratory, 08 2017. [p1]
- M. Tong, J. Song, and W. Chu. An Improved Algorithm of Congruent Matching Cells (CMC) Method for Firearm Evidence Identifications. *Journal of Research of the National Institute of Standards and Technology*, 120:102, Apr. 2015. ISSN 2165-7254. doi: 10.6028/jres.120.008. URL <https://nvlpubs.nist.gov/nistpubs/jres/120/jres.120.008.pdf>. [p2, 5, 6]
- X. A. Zheng, J. A. Soons, and R. M. Thompson. Nist ballistics toolmark research database, 2016. [p2]

Joseph Zemmels
Iowa State University Department of Statistics
Address
Unites States
(ORCID if desired)
jzemmels@iastate.edu

Heike Hofmann
Iowa State University Department of Statistics
Address
United States
(ORCID if desired)
hofmann@iastate.edu

Susan VanderPlas
University of Nebraska - Lincoln Department of Statistics
Address
United States
(ORCID if desired)
author3@work